

```

// A is printed before F
Semaphore Amutex = new Semaphore(1);
Semaphore Fmutex = new Semaphore(0);
thread P:{ 
    Amutex.acquire();
    print('A');
    Fmutex.release();
    Amutex.release();
    print('B');
    print('C');
}

thread Q:{ 
    print('E');
    Fmutex.acquire();
    print('F');
    print('G');
}

// using Semaphore to have an output of RACES
Semaphore Rmutex = new Semaphore(1);
Semaphore Amutex = new Semaphore(0);
Semaphore Emutex = new Semaphore(0);
thread P:{ 
    Amutex.acquire();
    print('A');
    print('C');
    Emutex.release();
}

thread Q:{ 
    Rmutex.acquire();
    print('R');
    Amutex.release();
    Rmutex.release();
    Emutex.acquire();
    print('E');
    print('S');
}

// using Semaphore to have an output of R / F / O / K / OK
Semaphore Imutex = new Semaphore(0);
Semaphore Omutex = new Semaphore(0);
Semaphore OKmutex = new Semaphore(0);
thread P:{ 
    print('R');
    Imutex.release();
    OKmutex.acquire();
    print("OK");
}

thread Q:{ 
    Imutex.acquire();
    print('I');
    Omutex.release();
    OKmutex.acquire();
    print("OK");
}

thread X:{ 
    Omutex.acquire();
    print('O');
    OKmutex.release();
    OKmutex.release();
    print("OK");
}

```

Readers/Writers

```

1  public void StartWrite() {
2      if (writers != 0 or readers != 0) { 
3          OKtoWrite.wait();
4      }
5      writers = writers + 1;
6  }
7
8  public void EndWrite() { 
9      writers = writers - 1;
10     if (OKtoRead.empty()) { 
11         OKtoWrite.signal();
12     } else { 
13         OKtoRead.signal();
14     }
15 }
16

```

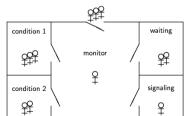
Readers/Writers

```

1 monitor RW {
2     int readers = 0;
3     int writers = 0;
4     condition OKtoRead, OKtoWrite;
5
6     public void StartRead() {
7         if (writers != 0 or not OKtoWrite.empty()) { 
8             OKtoRead.wait();
9         }
10        readers = readers + 1;
11        OKtoRead.signal();
12    }
13
14    public void EndRead() { 
15        readers = readers - 1;
16        if (readers==0) { 
17            OKtoWrite.signal();
18        }
19    }
20

```

signal and Continue: E = W < S (Java)



- ▶ Process from S which executes the signal continues execution
- ▶ Process from W which is unblocked joins competition for the lock
- ▶ Problem: signaling process can modify the condition after it executed the signal

```

//A hotel serves coffee using a dispenser.
//The dispenser has two buttons, so it can serve two clients at the same time.
//Every now and then the dispenser is refilled by an employee.
//There are multiple employees but only one should refill the
//dispenser at any given time. Moreover, she must ensure there are no
//clients when she refills the dispenser.
//Model this scenario using semaphores. Model both the client and the employee.
Semaphore mutex = new Semaphore(1);
Semaphore button = new Semaphore(2);

```

```

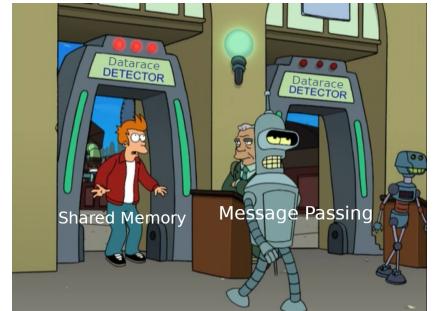
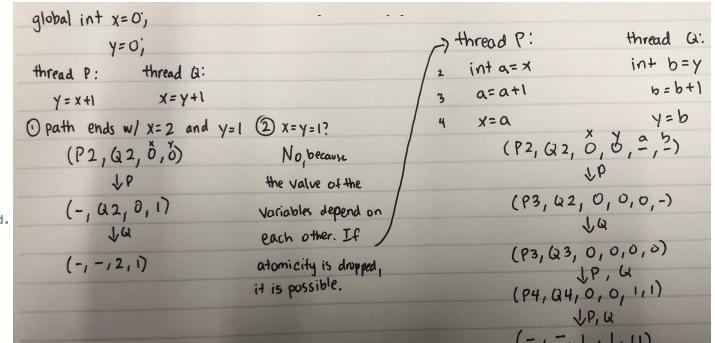
thread Client: {
    mutex.acquire();
    button.release();
}

thread Employee: {
    mutex.acquire();
    button.acquire();
    button.acquire();
    refillDispenser();
    button.release();
    button.release();
    mutex.release();
}

```

Livelock – loops and never hits CS

Deadlock – race condition in which two threads need two resources that the other



```

monitor Stack: {
    private Stack = new ArrayList();
    int size = 0;
    condition canPop;
    function push(i) {
        arr[size] = i;
        size++;
        canPop.signal();
    }
    function pop(){
        if(size == 0){
            canPop.wait();
        }
        size--;
        return arr[size+1];
    }
}

```

// We would like to model a control system for an automatic car wash. Each car traverses three phases: blast, rinse and dry. Each of these phases is executed by a machine. All vehicles follow these three phases in that exact order.

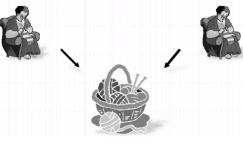
```

Semaphore[] permToProcess = { new Semaphore(0), new Semaphore(0), new Semaphore(0) };
Semaphore[] doneProcessing = { new Semaphore(0), new Semaphore(0), new Semaphore(0) };
Semaphore[] permToEnter = { new Semaphore(1), new Semaphore(1), new Semaphore(1) };

thread Car: {
    permToEnter[0].acquire(); // See if it is ok to enter station 0
    advanceToNextStation(); // assuming no cars are in any station, and the first car is in a station before station 0 that does not
    permToProcess[0].release(); // Start processing on car in station 0
    doneProcessing[0].acquire(); // Wait for the station thread to be done on station 0
    advanceToEnter[1].acquire(); // See if it is ok to enter station 1
    advanceToNextStation(); // Advance to station 1
    permToEnter[1].release(); // Allow cars into station 0
    permToProcess[1].acquire(); // Start processing on car in station 1
    doneProcessing[1].acquire(); // Wait for the station thread to be done on station 1
    advanceToEnter[2].acquire(); // See if it is ok to enter station 2
    advanceToNextStation(); // Advance to station 2
    permToEnter[2].release(); // Allow cars into station 2
}

thread MachineAtStation (i): {
    while (true) {
        permToProcess[i].acquire();
        // process vehicle
        doneProcessing[i].release();
    }
}

```



The Mutual Exclusion Problem (MEP)

The Mutual Exclusion Problem

Guarantee that:

1. **Mutex:** At any point in time, there is at most one thread in the critical section
2. **Absence of livelock:** If various threads try to enter the critical section, at least one of them will succeed
3. **Free from starvation:** A thread trying to enter its critical section will eventually be able to do so

Properties of Paths

- A property is an assertion that is true for every possible path
- Safety: A path never reaches a "bad" state
- Liveness: Eventually, every path will reach a "good" state

Synchronization

Mechanism that restricts the possible paths of a concurrent program in order to ensure safety and liveness properties.

Atomicity Assumption on Assignments

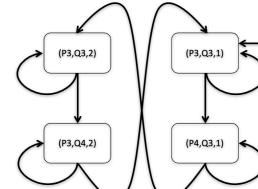
Atomic Operation

An operation is atomic if it is executed until completion without interleaving statements from other processes

- Atomic operations are the smallest unit in which a path can be listed
- We assume throughout this course that assignment statements are **atomic** for scalar values (i.e. numbers, booleans, chars, etc.)
- Eg. $x=2$ is atomic but $x=x+1$ is not.

Proving the Desired Properties I

- Mutex: Holds if all accessible states do not contain a state of the form $(p4, q4, turn)$ for some value of turn



Dekker's Algorithm (I + IV)

```
global boolean flag = false;
thread P: {
    // non-critical section
    while (flag) {
        // CRITICAL SECTION
        flag = true;
        // CRITICAL SECTION
        flag = false;
        // non-critical section
    }
}
while (cond) { is called a busy-wait loop
    Abbreviation:
        while (cond) 0 → await !cond
}
critical Section
```

Critical Section

A part of the program that accesses shared memory and which we wish to execute atomically

```
1 thread P: {
2     while(true) {
3         // non-critical section
4         entry to critical section;
5         // CRITICAL SECTION
6         exit from critical section;
7         // non-critical section
8     }
9 }
10 thread Q: {
11     while(true) {
12         // non-critical section
13         entry to critical section;
14         // CRITICAL SECTION
15         exit from critical section;
16         // non-critical section
17     }
18 }
```

Race Condition

Race Condition
Arises if two or more threads access the same variables or objects concurrently and at least one does updates

- Whether it happens depends on how threads are scheduled
- Once thread T1 starts doing something, it needs to "race" to finish it because if thread T2 looks at the shared variable before T1 is done, it may see something inconsistent
- Hard to detect

The Bakery Algorithm

```
global boolean[] choosing = new boolean[n]; // {false,...}
global int[] number = new int[n]; // {0,0,...0}

thread {
    // non-critical section
    choosing[id] = true;
    number[id] = 1 + maximum(number);
    choosing[id] = false;
    for (all other processes j) {
        await !choosing[j];
        await (number[j] == 0) or (number[id]<<number[j]);
    }
}

// CRITICAL SECTION
number[id] = 0;
// non-critical section
```

This algorithm solves the problem for n threads.

- Mutex: Yes
- Absence livelock: Yes
- Free from starvation: Yes

Right to insist on entering is passed between the two processes

Peterson's Algorithm (1981)

```
global int last = 1;
global boolean wantP = false;
global boolean wantQ = false;

thread P: {
    while (true) {
        // non-critical section
        wantP = true;
        last = 1;
        await !wantQ or last==2;
        // CRITICAL SECTION
        wantP = false;
        // non-critical section
    }
}
thread Q: {
    while (true) {
        // non-critical section
        wantQ = true;
        last = 2;
        await !wantP or last==1;
        // CRITICAL SECTION
        wantQ = false;
        // non-critical section
    }
}
```

Similar to Dekker except that if both want access, priority is given

Absence of livelock (AoL) means that if two or more processes are trying to enter their CS, then at least one must eventually succeed

- If our processes could **block** when trying to access the CS, then we would speak of **absence of deadlock**: there are no states that are not endstates and that have no outgoing edges
- If processes cannot block (as is our case for now) we talk about **absence of livelock** rather than AoD. Livelock means that even though processes are not blocked (i.e. in a waiting state) they make no progress (i.e., there is no path leading to the CS)

The Semaphore Solution for the MEP ($k = 1$)

#criticalSection: number of processes in their critical sections

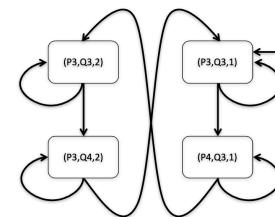
1. #criticalSection + permissions = 1
2. #criticalSection = #acquires - #releases

Item 1 guarantees:

- Mutual Exclusion ($\#criticalSection \leq 1$ since $0 \leq \text{permissions}$)
- Absence of deadlock (it never happens that $\text{permissions} = 0$ and $\#criticalSection = 0$)
- No starvation between two processes

Proving the Desired Properties II

- Absence of livelock (if two or more processes are trying to enter their CS, i.e. executing an await, then at least one must eventually succeed)
- Recall that the awaits are on line 3



Proving the Desired Properties III

Freedom from starvation (if any process is about to execute its protocol, then eventually it will succeed in entering the CS)

- Consider what happens if Q is trying to access its critical section but $turn = 1$ and P fails or loops in its NCS (*)

```
global int turn = 1;
global boolean wantP = false;
global boolean wantQ = false;

thread P: {
    while (true) {
        // non-critical section
        wantP = true;
        while (true) {
            // non-critical section
            wantP = true;
            while (true) {
                if (turn == 2) {
                    wantP = false;
                    await (turn++=1);
                    wantP = true;
                }
            }
            // CRITICAL SECTION
            turn = 2;
            wantP = false;
            wantQ = false;
            wantP = true;
        }
    }
}
thread Q: {
    while (true) {
        // non-critical section
        wantQ = true;
        await (turn++=1);
        // CRITICAL SECTION
        turn = 1;
        wantP = false;
        wantQ = false;
        wantP = true;
    }
}
```

Attempt III – Assessment

```
global boolean wantP = false;
global boolean wantQ = false;

thread P: {
    while (true) {
        // non-critical section
        wantP = true;
        await !wantQ;
        // CRITICAL SECTION
        wantP = false;
        // non-critical section
    }
}
thread Q: {
    while (true) {
        // non-critical section
        wantQ = true;
        await !wantP;
        // CRITICAL SECTION
        wantQ = false;
        // non-critical section
    }
}
```

Processes should:

1. either back out if they discover they are contending with the other (Naive back-out);
2. take turns (Dekker's algorithm);
3. give priority to the first one that wanted access (Peterson's algorithm).

Readers/Writers

- There are shared resources between two types of threads

- Readers: access the resource without modifying it
- Writers: access the resource and may modify it

Mutual exclusion is too restrictive

- Readers: can access simultaneously
- Writers: at most one at any given time

non-atomic	→ write is as assembly
- write a temp (int a = x)	
- store a stamp (a = a + 1)	
- set the variable to the temp value (x = a)	
atomic	→ do it

Condition Variables

Cond.wait()

- Always blocks the process and places it in the waiting queue of the variable Cond.
- When it blocks, it releases the mutex on the monitor.

Cond.signal()

- Unblocks the first process in the waiting queue of the variable Cond and sets it to the READY state
- If there are no processes in the waiting queue, it has no effect.

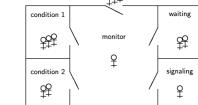
Cond.empty()

- Checks if waiting queue of Cond is empty or not

Example: Buffer of Size 1

```
1 monitor Buffer {
2     private Object buffer = null; // shared buffer
3     private condition full; // wait until space available
4     private condition empty; // wait until buffer available
5
6     public Object consume() {
7         if (buffer == null)
8             full.wait();
9         Object aux = buffer;
10        buffer = null;
11        empty.signal();
12        return aux;
13    }
14
15    public void produce(Object o) {
16        if (buffer != null)
17            empty.wait();
18        buffer = o;
19        full.signal();
20    }
}
```

Signal



► Two strategies:

- Signal and Urgent Wait: $E < S < W$ (classical monitor)
- Signal and Continue: $E = W < S$ (Java)

where the letters denote the precedence of

- S: signalling processes
- W: waiting processes
- E: processes blocked on entry

$E < S < W$

- When a process blocked on a condition variable is signaled, it immediately begins executing ahead of the signaling process
- Resumes at the instruction immediately following the call to wait that blocked it
- No need to check the condition
- Rationale: Signaling process changes the state of the monitor so that the condition now holds
- Cons: signaling process unnecessarily delayed (unless signal is the last operation)

Monitor that Defines a Semaphore

Dining Philosophers

```
1 monitor Semaphore {
2     private condition nonZero;
3     private int permissions;
4
5     public Semaphore(int n) {
6         this.permissions = n;
7     }
8
9     public void acquire() {
10        if (permissions == 0)
11            nonZero.wait();
12        permissions--;
13    }
14
15    public void release() {
16        permissions++;
17        nonZero.signal();
18    }
19
20    ...
}
```

ForkMonitor

```
1 monitor ForkMonitor {
2     int[] fork = {2,2,2,2,2};
3     condition[] OKtoEat; // 0-4
4
5     public void takeForks(integer i) {
6         if (fork[i] != 2) {
7             OKtoEat[i].wait();
8         }
9         fork[i+1] = fork[i+1] - 1;
10        fork[i-1] = fork[i-1] - 1;
11        ...
12
13        public void releaseForks(integer i) {
14            fork[i+1] = fork[i+1] + 1;
15            fork[i-1] = fork[i-1] + 1;
16            if (fork[i+1] == 2) {
17                OKtoEat[i+1].signal();
18            }
19            if (fork[i-1] == 2) {
20                OKtoEat[i-1].signal();
21            }
22        }
23    }
24}
```